

```
# import libraries
import pandas as pd #data processing, CSV file I/O
import numpy as np # linear algebra
import matplotlib.pyplot as plt #creates plot
import seaborn as sns
```

```
# For ignoring warning
import warnings
warnings.filterwarnings("ignore")
```

```
#Download csv file into our Notebook
#Default separator is comma
data=pd.read_csv("California_Houses.csv")
```

```
#Exploring and cleaning data
#look at the dataset within each district(block)
data
```

|                | Median_House_Value | Median_Income | Median_Age | Tot_Rooms |
|----------------|--------------------|---------------|------------|-----------|
| Tot_Bedrooms \ |                    |               |            |           |
| 0              | 452600.0           | 8.3252        | 41         | 880       |
| 129            |                    |               |            |           |
| 1              | 358500.0           | 8.3014        | 21         | 7099      |
| 1106           |                    |               |            |           |
| 2              | 352100.0           | 7.2574        | 52         | 1467      |
| 190            |                    |               |            |           |
| 3              | 341300.0           | 5.6431        | 52         | 1274      |
| 235            |                    |               |            |           |
| 4              | 342200.0           | 3.8462        | 52         | 1627      |
| 280            |                    |               |            |           |
| ...            | ...                | ...           | ...        | ...       |
| ...            |                    |               |            |           |
| 20635          | 78100.0            | 1.5603        | 25         | 1665      |
| 374            |                    |               |            |           |
| 20636          | 77100.0            | 2.5568        | 18         | 697       |
| 150            |                    |               |            |           |
| 20637          | 92300.0            | 1.7000        | 17         | 2254      |
| 485            |                    |               |            |           |
| 20638          | 84700.0            | 1.8672        | 18         | 1860      |
| 409            |                    |               |            |           |
| 20639          | 89400.0            | 2.3886        | 16         | 2785      |
| 616            |                    |               |            |           |

|   | Population | Households | Latitude | Longitude | Distance_to_coast |
|---|------------|------------|----------|-----------|-------------------|
| \ |            |            |          |           |                   |
| 0 | 322        | 126        | 37.88    | -122.23   | 9263.040773       |
| 1 | 2401       | 1138       | 37.86    | -122.22   | 10225.733072      |
| 2 | 496        | 177        | 37.85    | -122.24   | 8259.085109       |

|       |      |     |       |         |               |
|-------|------|-----|-------|---------|---------------|
| 3     | 558  | 219 | 37.85 | -122.25 | 7768.086571   |
| 4     | 565  | 259 | 37.85 | -122.25 | 7768.086571   |
| ...   | ...  | ... | ...   | ...     | ...           |
| 20635 | 845  | 330 | 39.48 | -121.09 | 162031.481121 |
| 20636 | 356  | 114 | 39.49 | -121.21 | 160445.433537 |
| 20637 | 1007 | 433 | 39.43 | -121.22 | 153754.341182 |
| 20638 | 741  | 349 | 39.43 | -121.32 | 152005.022239 |
| 20639 | 1387 | 530 | 39.37 | -121.24 | 146866.196892 |

|       | Distance_to_LA | Distance_to_SanDiego | Distance_to_SanJose \ |
|-------|----------------|----------------------|-----------------------|
| 0     | 556529.158342  | 735501.806984        | 67432.517001          |
| 1     | 554279.850069  | 733236.884360        | 65049.908574          |
| 2     | 554610.717069  | 733525.682937        | 64867.289833          |
| 3     | 555194.266086  | 734095.290744        | 65287.138412          |
| 4     | 555194.266086  | 734095.290744        | 65287.138412          |
| ...   | ...            | ...                  | ...                   |
| 20635 | 654530.186299  | 830631.543047        | 248510.058162         |
| 20636 | 659747.068444  | 836245.915229        | 246849.888948         |
| 20637 | 654042.214020  | 830699.573163        | 240172.220489         |
| 20638 | 657698.007703  | 834672.461887        | 238193.865909         |
| 20639 | 648723.337126  | 825569.179028        | 233282.769063         |

|       | Distance_to_SanFrancisco |
|-------|--------------------------|
| 0     | 21250.213767             |
| 1     | 20880.600400             |
| 2     | 18811.487450             |
| 3     | 18031.047568             |
| 4     | 18031.047568             |
| ...   | ...                      |
| 20635 | 222619.890417            |
| 20636 | 218314.424634            |
| 20637 | 212097.936232            |
| 20638 | 207923.199166            |
| 20639 | 205473.376575            |

[20640 rows x 14 columns]

*#No. of (Rows ,Columns)*

data.shape

(20640, 14)

```
#Checking for Duplicates
```

```
data.duplicated().sum()
```

```
0
```

```
#Preprocessing
```

```
#1-Check if data has any null values
```

```
data.info()
```

```
#if there is any nan-> not a number
```

```
#Drop null values and save: data.dropna(inplace=True)
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 20640 entries, 0 to 20639
```

```
Data columns (total 14 columns):
```

| #  | Column                   | Non-Null Count | Dtype   |
|----|--------------------------|----------------|---------|
| 0  | Median_House_Value       | 20640 non-null | float64 |
| 1  | Median_Income            | 20640 non-null | float64 |
| 2  | Median_Age               | 20640 non-null | int64   |
| 3  | Tot_Rooms                | 20640 non-null | int64   |
| 4  | Tot_Bedrooms             | 20640 non-null | int64   |
| 5  | Population               | 20640 non-null | int64   |
| 6  | Households               | 20640 non-null | int64   |
| 7  | Latitude                 | 20640 non-null | float64 |
| 8  | Longitude                | 20640 non-null | float64 |
| 9  | Distance_to_coast        | 20640 non-null | float64 |
| 10 | Distance_to_LA           | 20640 non-null | float64 |
| 11 | Distance_to_SanDiego     | 20640 non-null | float64 |
| 12 | Distance_to_SanJose      | 20640 non-null | float64 |
| 13 | Distance_to_SanFrancisco | 20640 non-null | float64 |

```
dtypes: float64(9), int64(5)
```

```
memory usage: 2.2 MB
```

```
#OR Checking for null values
```

```
data.isnull().sum()
```

|                      |   |
|----------------------|---|
| Median_House_Value   | 0 |
| Median_Income        | 0 |
| Median_Age           | 0 |
| Tot_Rooms            | 0 |
| Tot_Bedrooms         | 0 |
| Population           | 0 |
| Households           | 0 |
| Latitude             | 0 |
| Longitude            | 0 |
| Distance_to_coast    | 0 |
| Distance_to_LA       | 0 |
| Distance_to_SanDiego | 0 |
| Distance_to_SanJose  | 0 |

```
Distance_to_SanFrancisco    0
dtype: int64
```

```
#the dataset contains 20,640 samples and 14 features;
```

```
#all features are numerical features encoded as floating number or int
;
```

```
#there is no missing values.
```

```
#descriptive statistics of data
```

```
data.describe()
```

|             | Median_House_Value | Median_Income | Median_Age   |              |
|-------------|--------------------|---------------|--------------|--------------|
| Tot_Rooms \ |                    |               |              |              |
| count       | 20640.000000       | 20640.000000  | 20640.000000 | 20640.000000 |
| mean        | 206855.816909      | 3.870671      | 28.639486    | 2635.763081  |
| std         | 115395.615874      | 1.899822      | 12.585558    | 2181.615252  |
| min         | 14999.000000       | 0.499900      | 1.000000     | 2.000000     |
| 25%         | 119600.000000      | 2.563400      | 18.000000    | 1447.750000  |
| 50%         | 179700.000000      | 3.534800      | 29.000000    | 2127.000000  |
| 75%         | 264725.000000      | 4.743250      | 37.000000    | 3148.000000  |
| max         | 500001.000000      | 15.000100     | 52.000000    | 39320.000000 |

|             | Tot_Bedrooms | Population   | Households   | Latitude     |   |
|-------------|--------------|--------------|--------------|--------------|---|
| Longitude \ |              |              |              |              |   |
| count       | 20640.000000 | 20640.000000 | 20640.000000 | 20640.000000 |   |
| mean        | 537.898014   | 1425.476744  | 499.539680   | 35.631861    | - |
| std         | 421.247906   | 1132.462122  | 382.329753   | 2.135952     |   |
| min         | 1.000000     | 3.000000     | 1.000000     | 32.540000    | - |
| 25%         | 295.000000   | 787.000000   | 280.000000   | 33.930000    | - |
| 50%         | 435.000000   | 1166.000000  | 409.000000   | 34.260000    | - |
| 75%         | 647.000000   | 1725.000000  | 605.000000   | 37.710000    | - |
| max         | 6445.000000  | 35682.000000 | 6082.000000  | 41.950000    | - |

|       | Distance_to_coast | Distance_to_LA | Distance_to_SanDiego \ |
|-------|-------------------|----------------|------------------------|
| count | 20640.000000      | 2.064000e+04   | 2.064000e+04           |
| mean  | 40509.264883      | 2.694220e+05   | 3.981649e+05           |
| std   | 49140.039160      | 2.477324e+05   | 2.894006e+05           |
| min   | 120.676447        | 4.205891e+02   | 4.849180e+02           |
| 25%   | 9079.756762       | 3.211125e+04   | 1.594264e+05           |
| 50%   | 20522.019101      | 1.736675e+05   | 2.147398e+05           |
| 75%   | 49830.414479      | 5.271562e+05   | 7.057954e+05           |
| max   | 333804.686371     | 1.018260e+06   | 1.196919e+06           |

|       | Distance_to_SanJose | Distance_to_SanFrancisco |
|-------|---------------------|--------------------------|
| count | 20640.000000        | 20640.000000             |
| mean  | 349187.551219       | 386688.422291            |
| std   | 217149.875026       | 250122.192316            |
| min   | 569.448118          | 456.141313               |
| 25%   | 113119.928682       | 117395.477505            |
| 50%   | 459758.877000       | 526546.661701            |
| 75%   | 516946.490963       | 584552.007907            |
| max   | 836762.678210       | 903627.663298            |

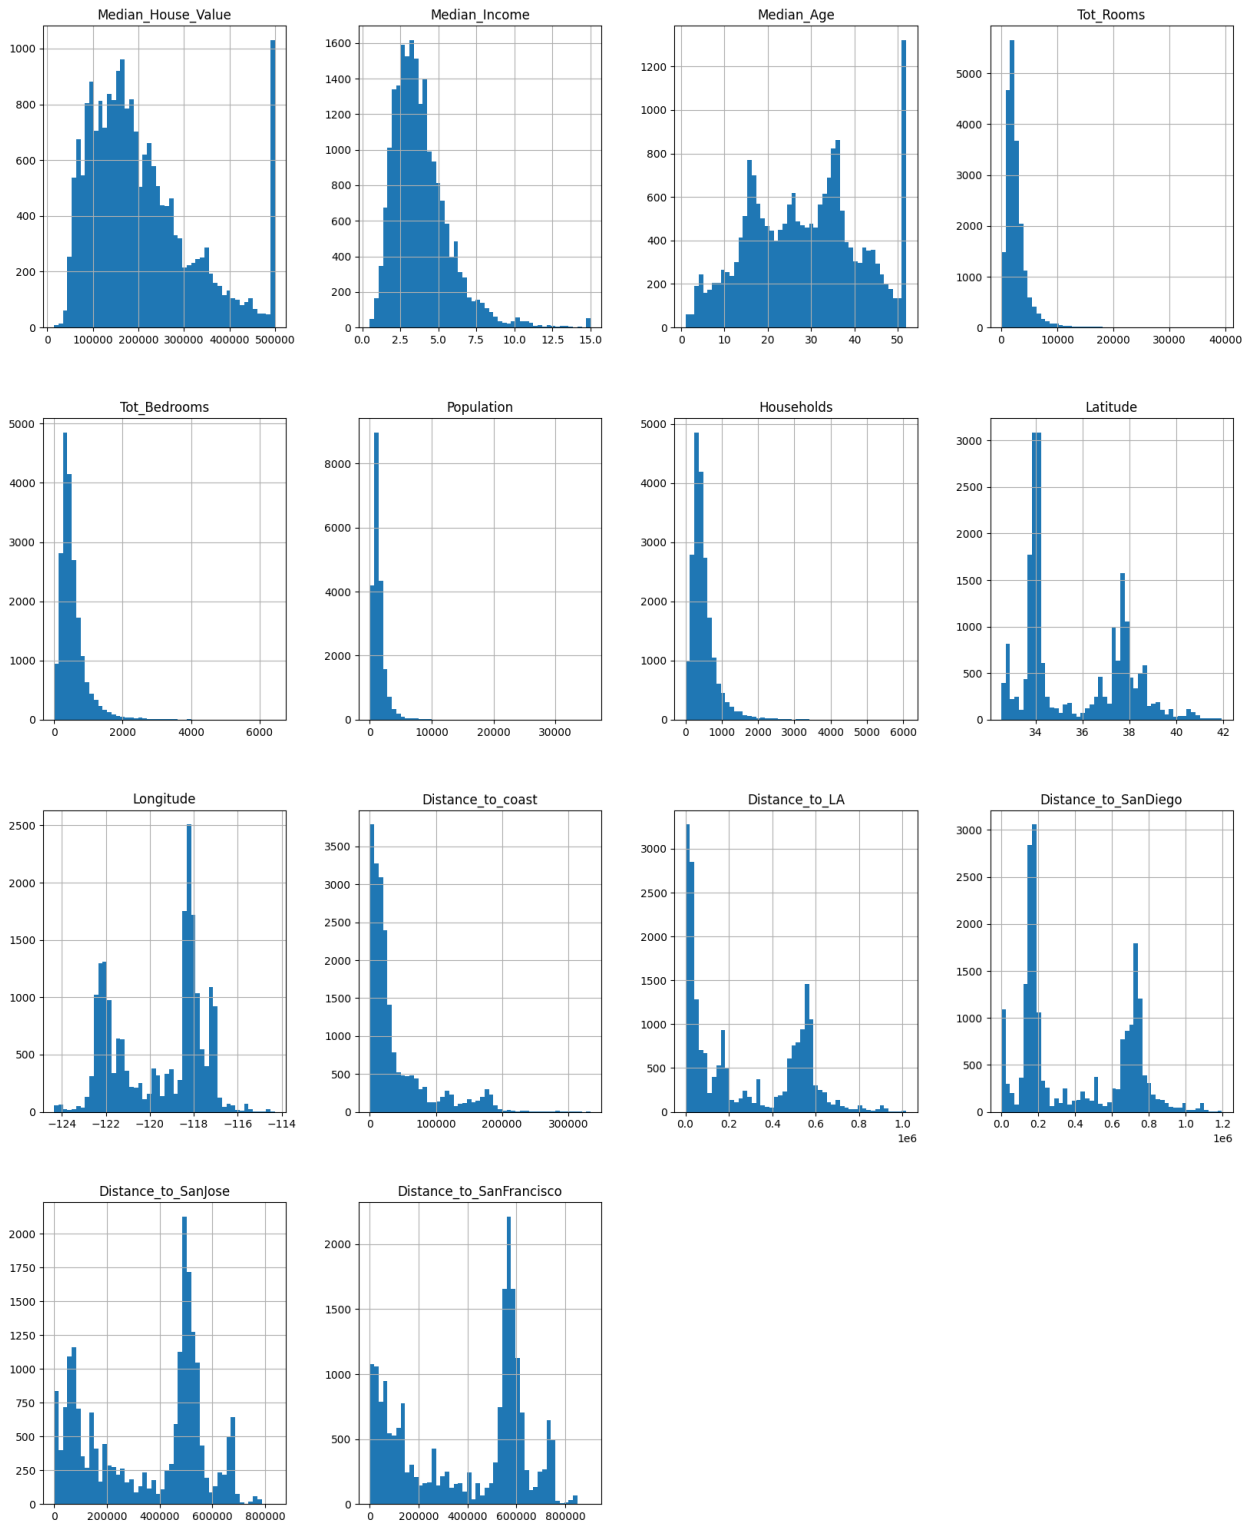
*#Let's check the distribution of the Target variable.*  
data['Median\_House\_Value'].value\_counts()

```
Median_House_Value
500001.0    965
137500.0    122
162500.0    117
112500.0    103
187500.0     93
...
359200.0     1
54900.0      1
377600.0     1
81200.0      1
47000.0      1
Name: count, Length: 3842, dtype: int64
```

data.hist(bins=50,figsize=(20,25))*#No of bins is No of chuncks to split data into*

```
array([[<Axes: title={'center': 'Median_House_Value'}>,
      <Axes: title={'center': 'Median_Income'}>,
      <Axes: title={'center': 'Median_Age'}>,
      <Axes: title={'center': 'Tot_Rooms'}>],
      [<Axes: title={'center': 'Tot_Bedrooms'}>,
      <Axes: title={'center': 'Population'}>,
      <Axes: title={'center': 'Households'}>,
      <Axes: title={'center': 'Latitude'}>],
      [<Axes: title={'center': 'Longitude'}>,
      <Axes: title={'center': 'Distance_to_coast'}>],
      ])
```

```
<Axes: title={'center': 'Distance_to_LA'}>,
<Axes: title={'center': 'Distance_to_SanDiego'}>],
[<Axes: title={'center': 'Distance_to_SanJose'}>,
<Axes: title={'center': 'Distance_to_SanFrancisco'}>, <Axes:
>,
<Axes: >]], dtype=object)
```

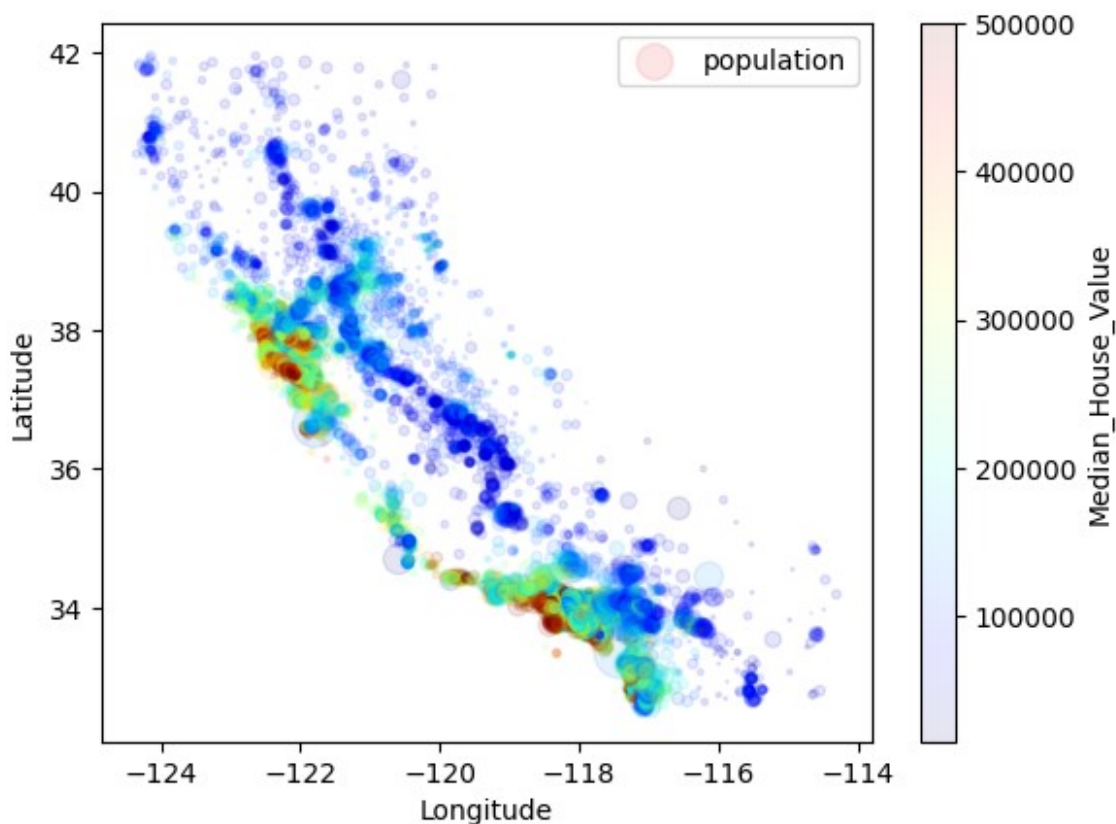


#Median house value has outliers at the very right where some houses that price much much higher than the rest of the population & they are all maxed out at 500K\$

```
#Means every house with at least 500k$ is labeled as 500k$  
#Also in MEdian house age all houses age above 50 years are grouped  
together
```

```
#plot the geographical features -> Map of California  
data.plot(kind="scatter",x="Longitude",y="Latitude",alpha=0.1,  
          s=data["Population"]/100,label="population",  
          c="Median_House_Value",cmap=plt.get_cmap("jet"))  
#alpha is used for transparency thus when points are laid over each  
other they are darker  
#darker points houses near cost are more expensive  
#sized of dots represents the population of the district  
#color corresponds to house value
```

```
<Axes: xlabel='Longitude', ylabel='Latitude'>
```



```
#create correlation matrix in a table corr_matrix->shows how closely  
related 2 variables are  
#or how much one variable changes as the other one changes  
#correlation of a variable with itself=1 highest possible value of  
correlation
```

```
corr_matrix=data.corr()  
#correlation Range [-1,+1]  
#corr=-1: exactly the opposite
```

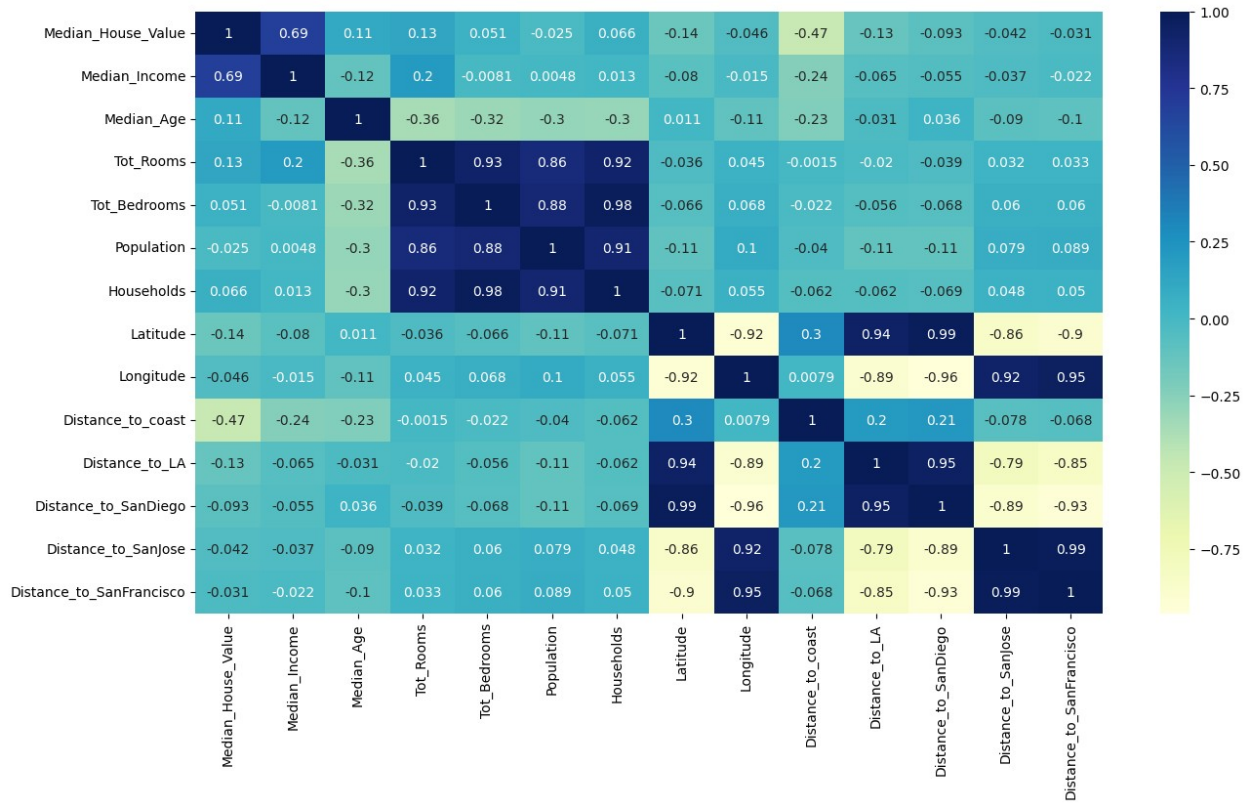


```
#corr=0: No relation
#corr1: exactly the same
# if close to one or -1 they are closely related/
corr_matrix["Median_House_Value"].sort_values(ascending=False)
#Specify column and sort correlations in descending order
```

```
Median_House_Value      1.000000
Median_Income            0.688075
Tot_Rooms                0.134153
Median_Age               0.105623
Households              0.065843
Tot_Bedrooms            0.050594
Population              -0.024650
Distance_to_SanFrancisco -0.030559
Distance_to_SanJose     -0.041590
Longitude                -0.045967
Distance_to_SanDiego    -0.092510
Distance_to_LA          -0.130678
Latitude                 -0.144160
Distance_to_coast       -0.469350
Name: Median_House_Value, dtype: float64
```


```
#plot heatmap of correlation
plt.figure(figsize=(15,8))
sns.heatmap(corr_matrix,annot=True, cmap="YlGnBu")
```

```
<Axes: >
```



```
#plot heatmap for correlated data above 50%
kot = corr_matrix[corr_matrix>=.50]
plt.figure(figsize=(12,8))
sns.heatmap(kot, cmap="Blues")
```

<Axes: >



```
#feature engineering: Combine features to generate new interesting features
#Exp: -Total no. of bedrooms
#      -Total no. of rooms
#New feature: The number of bedrooms per room
data['Bedroom_Ratio']=data['Tot_Bedrooms']/data['Tot_Rooms']
#New feature: How many rooms per household
# data['Household_Ratio']=data['Tot_Rooms']/data['Households']
# data['Rooms_Population']=data['Tot_Rooms']*data['Population']
# data['Bedrooms_Population']=data['Tot_Bedrooms']*data['Population']
# data['Bedrooms_Households']=data['Tot_Bedrooms']*data['Households']
# data['Population_Households']=data['Population']*data['Households']
data['Value_Income']=data['Median_House_Value']*data['Median_Income']
data['Latitude_LA']=data['Latitude']*data['Distance_to_LA']
data['Latitude_SanDiego']=data['Latitude']*data['Distance_to_SanDiego']
]
data['Longitude_SanJose']=data['Longitude']*data['Distance_to_SanJose']
]
data['Longitude_SanFrancisco']=data['Longitude']*data['Distance_to_San Francisco']
]
```

```
data['LA_SanDiego']=data['Distance_to_LA']*data['Distance_to_SanDiego']
data['SanJose_SanFrancisco']=data['Distance_to_SanJose']*data['Distance_to_SanFrancisco']
```

*#To remove skewness in data , +1 is used to avoid log(0)!!!*

```
data['Tot_Rooms']=np.log(data['Tot_Rooms']+1)
data['Tot_Bedrooms']=np.log(data['Tot_Bedrooms']+1)
data['Population']=np.log(data['Population']+1)
data['Households']=np.log(data['Households']+1)
data['Distance_to_coast']=np.log(data['Distance_to_coast']+1)
```

*#Creating features and labels dataset*

*#x is the features dataset without target(label), Drop the target ->Median\_House\_Value*

```
x=data.drop(['Median_House_Value'],axis =1 ) # axis=1 to drop column
```

*#y is only the target (label) dataset*

```
y=data['Median_House_Value']
```

x

|       | Median_Income | Median_Age | Tot_Rooms | Tot_Bedrooms | Population |
|-------|---------------|------------|-----------|--------------|------------|
| 0     | 8.3252        | 41         | 6.781058  | 4.867534     | 5.777652   |
| 1     | 8.3014        | 21         | 8.867850  | 7.009409     | 7.784057   |
| 2     | 7.2574        | 52         | 7.291656  | 5.252273     | 6.208590   |
| 3     | 5.6431        | 52         | 7.150701  | 5.463832     | 6.326149   |
| 4     | 3.8462        | 52         | 7.395108  | 5.638355     | 6.338594   |
| ...   | ...           | ...        | ...       | ...          | ...        |
| 20635 | 1.5603        | 25         | 7.418181  | 5.926926     | 6.740519   |
| 20636 | 2.5568        | 18         | 6.548219  | 5.017280     | 5.877736   |
| 20637 | 1.7000        | 17         | 7.720905  | 6.186209     | 6.915723   |
| 20638 | 1.8672        | 18         | 7.528869  | 6.016157     | 6.609349   |
| 20639 | 2.3886        | 16         | 7.932362  | 6.424869     | 7.235619   |

|                  | Households | Latitude | Longitude | Distance_to_coast |
|------------------|------------|----------|-----------|-------------------|
| Distance_to_LA \ |            |          |           |                   |
| 0                | 4.844187   | 37.88    | -122.23   | 9.133896          |
| 556529.158342    |            |          |           |                   |

|               |          |       |         |           |
|---------------|----------|-------|---------|-----------|
| 1             | 7.037906 | 37.86 | -122.22 | 9.232760  |
| 554279.850069 |          |       |         |           |
| 2             | 5.181784 | 37.85 | -122.24 | 9.019190  |
| 554610.717069 |          |       |         |           |
| 3             | 5.393628 | 37.85 | -122.25 | 8.957908  |
| 555194.266086 |          |       |         |           |
| 4             | 5.560682 | 37.85 | -122.25 | 8.957908  |
| 555194.266086 |          |       |         |           |
| ...           | ...      | ...   | ...     | ...       |
| ...           |          |       |         |           |
| 20635         | 5.802118 | 39.48 | -121.09 | 11.995552 |
| 654530.186299 |          |       |         |           |
| 20636         | 4.744932 | 39.49 | -121.21 | 11.985715 |
| 659747.068444 |          |       |         |           |
| 20637         | 6.073045 | 39.43 | -121.22 | 11.943118 |
| 654042.214020 |          |       |         |           |
| 20638         | 5.857933 | 39.43 | -121.32 | 11.931675 |
| 657698.007703 |          |       |         |           |
| 20639         | 6.274762 | 39.37 | -121.24 | 11.897284 |
| 648723.337126 |          |       |         |           |

|                 | Distance_to_SanJose | Distance_to_SanFrancisco |
|-----------------|---------------------|--------------------------|
| Badroom_Ratio \ |                     |                          |
| 0               | 67432.517001        | 21250.213767             |
| 0.146591        |                     |                          |
| 1               | 65049.908574        | 20880.600400             |
| 0.155797        |                     |                          |
| 2               | 64867.289833        | 18811.487450             |
| 0.129516        |                     |                          |
| 3               | 65287.138412        | 18031.047568             |
| 0.184458        |                     |                          |
| 4               | 65287.138412        | 18031.047568             |
| 0.172096        |                     |                          |

|          |               |               |     |
|----------|---------------|---------------|-----|
| ...      | ...           | ...           | ... |
| ...      |               |               |     |
| 20635    | 248510.058162 | 222619.890417 |     |
| 0.224625 |               |               |     |
| 20636    | 246849.888948 | 218314.424634 |     |
| 0.215208 |               |               |     |
| 20637    | 240172.220489 | 212097.936232 |     |
| 0.215173 |               |               |     |
| 20638    | 238193.865909 | 207923.199166 |     |
| 0.219892 |               |               |     |
| 20639    | 233282.769063 | 205473.376575 |     |
| 0.221185 |               |               |     |

|                     | Value_Income | Latitude_LA  | Latitude_SanDiego |
|---------------------|--------------|--------------|-------------------|
| Longitude_SanJose \ |              |              |                   |
| 0                   | 3767985.52   | 2.108132e+07 | 2.786081e+07      |

```

8.242277e+06
1      2976051.90  2.098504e+07      2.776035e+07  -
7.950400e+06
2      2555330.54  2.099202e+07      2.776395e+07  -
7.929378e+06
3      1925990.03  2.101410e+07      2.778551e+07  -
7.981353e+06
4      1316169.64  2.101410e+07      2.778551e+07  -
7.981353e+06
...      ...      ...      ...      ..
.
20635      121859.43  2.584085e+07      3.279333e+07  -
3.009208e+07
20636      197129.28  2.605341e+07      3.302335e+07  -
2.992068e+07
20637      156910.00  2.578888e+07      3.275448e+07  -
2.911368e+07
20638      158151.84  2.593303e+07      3.291114e+07  -
2.889768e+07
20639      213540.84  2.554024e+07      3.250266e+07  -
2.828320e+07

```

|       | Longitude_SanFrancisco | LA_SanDiego  | SanJose_SanFrancisco |
|-------|------------------------|--------------|----------------------|
| 0     | -2.597414e+06          | 4.093282e+11 | 1.432955e+09         |
| 1     | -2.552027e+06          | 4.064184e+11 | 1.358281e+09         |
| 2     | -2.299516e+06          | 4.068212e+11 | 1.220250e+09         |
| 3     | -2.204296e+06          | 4.075655e+11 | 1.177195e+09         |
| 4     | -2.204296e+06          | 4.075655e+11 | 1.177195e+09         |
| ...   | ...                    | ...          | ...                  |
| 20635 | -2.695704e+07          | 5.436734e+11 | 5.532328e+10         |
| 20636 | -2.646189e+07          | 5.517108e+11 | 5.389089e+10         |
| 20637 | -2.571051e+07          | 5.433126e+11 | 5.094003e+10         |
| 20638 | -2.522524e+07          | 5.489624e+11 | 4.952603e+10         |
| 20639 | -2.491159e+07          | 5.355660e+11 | 4.793340e+10         |

[20640 rows x 21 columns]

y

```

0      452600.0
1      358500.0
2      352100.0
3      341300.0
4      342200.0
...
20635      78100.0
20636      77100.0
20637      92300.0
20638      84700.0

```

```
20639      89400.0
Name: Median_House_Value, Length: 20640, dtype: float64
```

```
from sklearn.model_selection import train_test_split

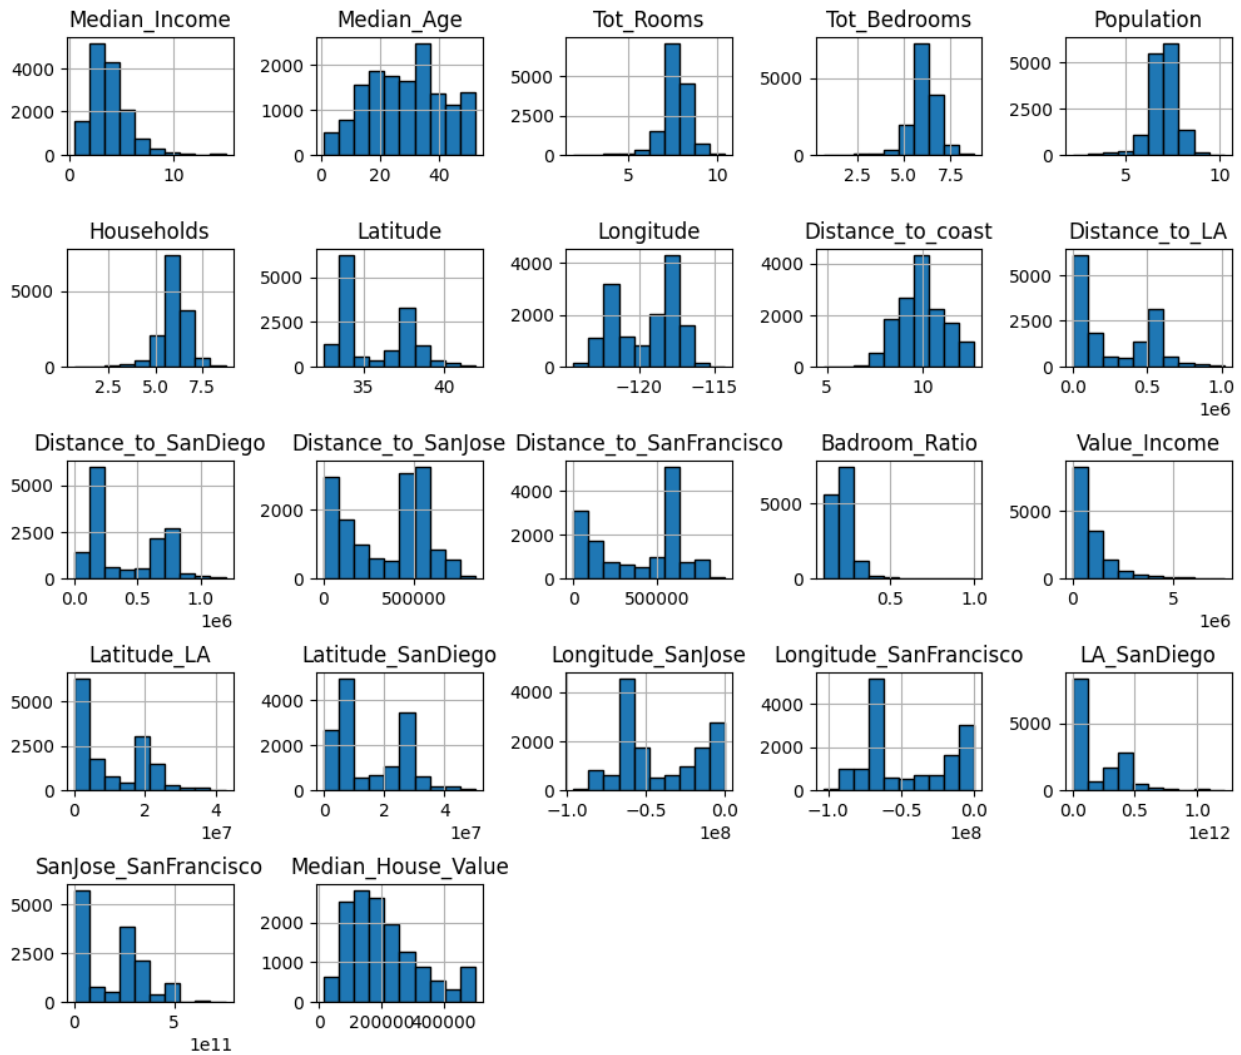
# set aside 15% of train and test data for evaluation
x_train, x_test, y_train, y_test = train_test_split(x, y,
    test_size=0.3, shuffle = True, random_state = 1900)
# Use the same function above for the validation set
x_test, x_val, y_test, y_val = train_test_split(x_test, y_test,
    test_size=0.5, random_state= 8) # 0.3 x 0.5= 0.15

#Data Splitting
print("x_train shape: {}".format(x_train.shape))
print("y_train shape: {}".format(y_train.shape))
print("x_val shape: {}".format(x_val.shape))
print("y_val shape: {}".format(y_val.shape))
print("x_test shape: {}".format(x_test.shape))
print("y_test shape: {}".format(y_test.shape))

x_train shape: (14448, 21)
y_train shape: (14448,)
x_val shape: (3096, 21)
y_val shape: (3096,)
x_test shape: (3096, 21)
y_test shape: (3096,)


#use only train data
#to get again the combined data frame but only for the training data
train_data=x_train.join(y_train)

# Data Exploration
#To get a histogram of individual features in Training set
train_data.hist(figsize=(12,10),edgecolor="black")
plt.subplots_adjust(hspace=0.8, wspace=0.5)
```



```
plt.figure(figsize=(15,8))
sns.scatterplot(
    data=train_data, x="Latitude",y="Longitude",
    hue="Median_House_Value" ,palette="viridis",
    alpha=0.5
)
plt.legend(title="MedHouseVal", bbox_to_anchor=(1.05, 0.95),
loc="upper left")
_ = plt.title("Median house value depending of\n their spatial
location")
```





```
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
#linear regression is called ordinary least squares-> OLS
OLS= Pipeline([
#Scaling input No need to scale output
    ('scaler', StandardScaler()),
    ('regressor', LinearRegression())
])
OLS.fit(x_train,y_train)

Pipeline(steps=[('scaler', StandardScaler()),
                 ('regressor', LinearRegression())])

# Display intercept and coefficients of the OLS model and R^2
print("Intercept is " + str(OLS['regressor'].intercept_))
print("The set of coefficients are " + str(OLS['regressor'].coef_))
#R^2 is the measure of performance of the OLS regression
print("The R-Squared value is " + str(OLS.score (x_train,y_train)))
```

```
Intercept is 207809.14168050993
The set of coefficients are [ -58070.26907658    -4782.89824866
 71598.85914909   -51330.44625947
 -40704.57555972    24227.22027286   -315641.77479945
 38271.41081192
 -13564.84107437   -145035.58914444  -1041973.07345917
 2894705.00070248
 -1210452.59847144    8928.02964488   128176.14191814
 64062.65017811
 1288718.83829814   2819391.9075079   -1160209.33682804
```

62609.68754738

-157230.57494066]

The R-Squared value is 0.848097947453569

```
#getting the best Hyperparameter for Ridge regression
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import Ridge
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
# set aside 15% of train and test data for evaluation
x_train, x_test, y_train, y_test = train_test_split(x, y,
    test_size=0.3, shuffle = True, random_state = 1900)
# Use the same function above for the validation set
x_test, x_val, y_test, y_val = train_test_split(x_test, y_test,
    test_size=0.5, random_state= 8) # 0.3 x 0.5= 0.15
errors =[]
#we will use our validation set
values = [1e-15, 1e-10, 1e-18, 1e-13, 1e-12, 1, 5, 10, 20, 30, 35, 40,
45, 50, 55, 100]
for i in range(len(values)):
    ridge= Pipeline([
#Scaling input No need to scale output
        ('scaler', StandardScaler()),
        ('regressor', Ridge(alpha=values[i]))
    ])
    ridge.fit(x_val,y_val)
    prediction = ridge.predict(x_val)
    MSE=mean_squared_error(y_val,prediction)
    errors.append(MSE)
index_of_best_alpha = np.array(errors).argmin()
best_alpha = values[index_of_best_alpha]
print("the best hyperparameter of Ridge regression: "+str(best_alpha))

ridge= Pipeline([
#Scaling input No need to scale output
    ('scaler', StandardScaler()),
    ('ridge_regressor', Ridge(alpha=best_alpha))
])
ridge.fit(x_train,y_train)
# print(ridge.score(x_train,y_train))

the best hyperparameter of Ridge regression: 1e-15

Pipeline(steps=[('scaler', StandardScaler()),
                ('ridge_regressor', Ridge(alpha=1e-15))])

#applying Ridge regression with another way to get best hyperparameter
from sklearn.linear_model import Ridge
```

```

from sklearn.model_selection import GridSearchCV

parameters = {'alpha':[1e-15,1e-10,1e-18,1e-13,1e-12,1,5,10,20,30,35,40,45,50,55,100]}
#chooseing best Hyperparameter by GridSearchCV
ridge_regressor=GridSearchCV(Ridge(),parameters,cv=5)
# ridge_regressor.fit(x_train,y_train)
ridge= Pipeline([
#Scaling input No need to scale output
    ('scaler', StandardScaler()),
    ('ridge_regressor', ridge_regressor)
])
#compare the Hyperparameter values between the 2 diffrent methods
ridge.fit(x_train,y_train)
if(ridge['ridge_regressor'].best_params_['alpha']==best_alpha):
    print("As we can see the Hyperparameter of the two methods have the same value which is: "+str(best_alpha))
else:
    print("As we can see the Hyperparameter of the two methods are different \n method-1= "+str(best_alpha)+"\n method-2= "+str(ridge['ridge_regressor'].best_params_['alpha']))

print(" How ever the diffrence are very small and both of them got the same score: "+str(ridge.score(x_train,y_train)))

```

As we can see the Hyperparameter of the two methods are different

method-1= 1e-15

method-2= 1e-10

How ever the diffrence are very small and both of them got the same score: 0.848097947453569

*#getting the best Hyperparameter for Lasso regression*

*#we will use our validation setonly this time because it is the requierd method*

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.metrics import mean_squared_error
```

```
from sklearn.linear_model import Lasso
```

```
errors =[]
```

```
values = [1e-15, 1e-10, 1e-18, 1e-13, 1e-12, 1, 5, 10, 20, 30, 35, 40, 45, 50, 55, 100]
```

```
for i in range(len(values)):
```

```
    lasso = Lasso(alpha=values[i])
```

```
    lasso.fit(x_val,y_val)
```

```
    prediction = lasso.predict(x_val)
```

```
    MSE=mean_squared_error(y_val,prediction)
```

```
    errors.append(MSE)
```

```
index_of_best_alpha = np.array(errors).argmin()
```

```
best_alpha_l = values[index_of_best_alpha]
```

```

print("the best hyperparameter of lasso regression:
"+str(best_alpha_l))
lasso= Pipeline([
#Scaling input No need to scale output
    ('scaler', StandardScaler()),
    ('ridge_regressor', Lasso(alpha=best_alpha_l))
])
lasso.fit(x_train,y_train)

the best hyperparameter of lasso regression: 1e-12
Pipeline(steps=[('scaler', StandardScaler()),
                 ('ridge_regressor', Lasso(alpha=1e-12))])

from sklearn.metrics import mean_squared_error
#predicting with OLS
y_pred=OLS.predict(x_val)#The predicted values based on validate
dataset
MSE=mean_squared_error(y_val,y_pred)
MSE
print("Mean square error: " + str( MSE))

Mean square error: 1844847809.5235841

y_pred
array([248695.33350802,  84759.68034878, 156303.63855013, ...,
       184375.74545651, 233194.41810921, 459402.60582285])

RMSE= np.sqrt(MSE)
print("Root mean square error: " + str( RMSE))

Root mean square error: 42951.69157930319

from sklearn.metrics import mean_absolute_error as MAE
# calculate MAE
error = MAE(y_val,y_pred)

# display
print("Mean absolute error : " + str(error))

Mean absolute error : 29667.89139684592

y_pred_R=ridge.predict(x_val)#The predicted values based on validate
dataset
#predicting with OLS
MSE_R=mean_squared_error(y_val,y_pred_R)
MSE_R
print("Mean square error for Ridge Regression: " + str( MSE_R))

Mean square error for Ridge Regression: 1844847809.4254832

```

```

RMSE_R= np.sqrt(MSE_R)
print("Root mean square error for Ridge Regression: " + str( RMSE_R))

Root mean square error for Ridge Regression: 42951.6915781612

from sklearn.metrics import mean_absolute_error as MAE
# calculate MAE
error = MAE(y_val,y_pred_R)

# display
print("Mean absolute error : " + str(error))

Mean absolute error : 29667.891399970013


y_pred_L=lasso.predict(x_val)#The predicted values based on validate
dataset
#predicting with OLS
MSE_L=mean_squared_error(y_val,y_pred_L)
MSE_L
print("Mean square error for Lasso Regression: " + str( MSE_L))

Mean square error for Lasso Regression: 1903693311.5767913

RMSE_L= np.sqrt(MSE_L)
print("Root mean square error for Lasso Regression: " + str( RMSE_L))

Root mean square error for Lasso Regression: 43631.3340568082

from sklearn.metrics import mean_absolute_error as MAE
# calculate MAE
error = MAE(y_val,y_pred_L)

# display
print("Mean absolute error for Lasso regression : " + str(error))

Mean absolute error for Lasso regression : 30304.374774314463

Performance=pd.DataFrame({'PREDICTIONS':y_pred,'ACTUAL VALUES':y_val})
#creating new column error
Performance['Error']=Performance['ACTUAL VALUES']-
Performance['PREDICTIONS']
Performance.head()

```

|       | PREDICTIONS   | ACTUAL VALUES | Error         |
|-------|---------------|---------------|---------------|
| 17909 | 248695.333508 | 247600.0      | -1095.333508  |
| 2451  | 84759.680349  | 67000.0       | -17759.680349 |
| 15005 | 156303.638550 | 133700.0      | -22603.638550 |
| 9346  | 275307.111164 | 316300.0      | 40992.888836  |
| 3543  | 350063.426898 | 357100.0      | 7036.573102   |

```
#preparing data for plotting :
Performance.reset_index(drop=True,inplace=True)
#new column index-> No. of observations
Performance.reset_index(inplace=True)
Performance.head()
```

|   | index | PREDICTIONS   | ACTUAL VALUES | Error         |
|---|-------|---------------|---------------|---------------|
| 0 | 0     | 248695.333508 | 247600.0      | -1095.333508  |
| 1 | 1     | 84759.680349  | 67000.0       | -17759.680349 |
| 2 | 2     | 156303.638550 | 133700.0      | -22603.638550 |
| 3 | 3     | 275307.111164 | 316300.0      | 40992.888836  |
| 4 | 4     | 350063.426898 | 357100.0      | 7036.573102   |

```
Performance_l=pd.DataFrame({'PREDICTIONS_l':y_pred_L,'ACTUAL
VALUES_l':y_val})
#creating new column error
Performance_l['Error_l']=Performance_l['ACTUAL VALUES_l']-
Performance_l['PREDICTIONS_l']
```

```
Performance_r=pd.DataFrame({'PREDICTIONS_r':y_pred_R,'ACTUAL
VALUES_r':y_val})
#creating new column error
Performance_r['Error_l']=Performance_r['ACTUAL VALUES_r']-
Performance_r['PREDICTIONS_r']
```


```
#preparing data for plotting :
Performance_l.reset_index(drop=True,inplace=True)
#new column index-> No. of observations
Performance_l.reset_index(inplace=True)
Performance_l.head()
```

|   | index | PREDICTIONS_l | ACTUAL VALUES_l | Error_l       |
|---|-------|---------------|-----------------|---------------|
| 0 | 0     | 244352.327026 | 247600.0        | 3247.672974   |
| 1 | 1     | 95647.921289  | 67000.0         | -28647.921289 |
| 2 | 2     | 154470.592652 | 133700.0        | -20770.592652 |
| 3 | 3     | 277081.777510 | 316300.0        | 39218.222490  |
| 4 | 4     | 349393.830291 | 357100.0        | 7706.169709   |

```
#preparing data for plotting :
Performance_r.reset_index(drop=True,inplace=True)
#new column index-> No. of observations
Performance_r.reset_index(inplace=True)
Performance_r.head()
```

|   | index | PREDICTIONS_r | ACTUAL VALUES_r | Error_l       |
|---|-------|---------------|-----------------|---------------|
| 0 | 0     | 248695.333571 | 247600.0        | -1095.333571  |
| 1 | 1     | 84759.680540  | 67000.0         | -17759.680540 |
| 2 | 2     | 156303.638538 | 133700.0        | -22603.638538 |
| 3 | 3     | 275307.111338 | 316300.0        | 40992.888662  |
| 4 | 4     | 350063.426986 | 357100.0        | 7036.573014   |

```
#plot the residuals(errors)
fig=plt.figure(figsize=(10,15))
#plot bar chart x-axis-> index , y-axis->error
plt.bar('index','Error',data=Performance[:50],color='black',width=0.3)
#data=Performance[:50] this means getting data from the performance
dataset's first 50 observaions to be clear
plt.xlabel("Observations")
plt.ylabel("Residuals")
plt.show()
#Error is positive means underestimating: Prediction lesser than the
actual value
#Error is negative means overestimating: Prediction higher than the
actual value
```





```
import statsmodels.api as sm #specialized library for statisticals
model Displays results better than sklearn
x_train=sm.add_constant(x_train)
#add constant =1 o train data set this important to get statsmodel
linear regression with intercept correctly identical to sklearn
x_train.head()
```

|              | const | Median_Income | Median_Age | Tot_Rooms | Tot_Bedrooms |
|--------------|-------|---------------|------------|-----------|--------------|
| Population \ |       |               |            |           |              |
| 6791         | 1.0   | 3.1250        | 44         | 6.960348  | 5.529429     |
| 6.848005     |       |               |            |           |              |
| 2888         | 1.0   | 1.8319        | 36         | 7.271704  | 5.855072     |
| 6.961296     |       |               |            |           |              |
| 3402         | 1.0   | 6.2037        | 32         | 6.405228  | 4.691348     |
| 5.752573     |       |               |            |           |              |
| 9330         | 1.0   | 6.1359        | 16         | 4.615121  | 3.044522     |
| 3.828641     |       |               |            |           |              |
| 12520        | 1.0   | 1.1458        | 48         | 6.995766  | 6.001415     |
| 6.831954     |       |               |            |           |              |

|       | Households | Latitude | Longitude | Distance_to_coast | ... | \ |
|-------|------------|----------|-----------|-------------------|-----|---|
| 6791  | 5.549076   | 34.08    | -118.15   | 10.413730         | ... |   |
| 2888  | 5.834811   | 35.39    | -118.99   | 11.710204         | ... |   |
| 3402  | 4.736198   | 34.27    | -118.35   | 10.359508         | ... |   |
| 9330  | 3.258097   | 37.96    | -122.50   | 6.738164          | ... |   |
| 12520 | 5.820083   | 38.55    | -121.47   | 10.888868         | ... |   |

|       | Distance_to_SanJose | Distance_to_SanFrancisco | Bedroom_Ratio | \ |
|-------|---------------------|--------------------------|---------------|---|
| 6791  | 495140.641637       | 563175.694111            | 0.238367      |   |
| 2888  | 338152.763741       | 405760.086533            | 0.242003      |   |
| 3402  | 467129.274368       | 535164.210208            | 0.178808      |   |
| 9330  | 87624.405549        | 21546.591499             | 0.200000      |   |
| 12520 | 140050.130972       | 120454.322106            | 0.369386      |   |

|                     | Value_Income | Latitude_LA  | Latitude_SanDiego |   |
|---------------------|--------------|--------------|-------------------|---|
| Longitude_SanJose \ |              |              |                   |   |
| 6791                | 642500.00    | 3.123455e+05 | 6.042801e+06      | - |
| 5.850087e+07        |              |              |                   |   |
| 2888                | 101487.26    | 5.791362e+06 | 1.209503e+07      | - |
| 4.023680e+07        |              |              |                   |   |
| 3402                | 1274239.98   | 8.949827e+05 | 7.025104e+06      | - |
| 5.528475e+07        |              |              |                   |   |
| 9330                | 1303878.75   | 2.197829e+07 | 2.876226e+07      | - |
| 1.073399e+07        |              |              |                   |   |
| 12520               | 74935.32     | 2.226577e+07 | 2.915958e+07      | - |
| 1.701189e+07        |              |              |                   |   |

|      | Longitude_SanFrancisco | LA_SanDiego  | SanJose_SanFrancisco |
|------|------------------------|--------------|----------------------|
| 6791 | -6.653921e+07          | 1.625078e+09 | 2.788512e+11         |
| 2888 | -4.828139e+07          | 5.592763e+10 | 1.372089e+11         |

|       |               |              |              |
|-------|---------------|--------------|--------------|
| 3402  | -6.333668e+07 | 5.353517e+09 | 2.499909e+11 |
| 9330  | -2.639457e+06 | 4.386969e+11 | 1.888007e+09 |
| 12520 | -1.463159e+07 | 4.368880e+11 | 1.686964e+10 |

[5 rows x 22 columns]

```
#sm.OLS(label dataset, features dataset that includes constant value).fit
```

```
nicerOLS=sm.OLS(y_train,x_train).fit()
```

```
nicerOLS.summary()
```

```
#comment: Constant is now the intercept
```

```
<class 'statsmodels.iolib.summary.Summary'>
```

```
"""
```

## OLS Regression Results

```
=====
=====
```

```
Dep. Variable:      Median_House_Value    R-squared:
```

```
0.848
```

```
Model:              OLS    Adj. R-squared:
```

```
0.848
```

```
Method:              Least Squares    F-statistic:
```

```
4026.
```

```
Date:              Wed, 01 Nov 2023    Prob (F-statistic):
```

```
0.00
```

```
Time:              23:21:17    Log-Likelihood:      -
```

```
1.7535e+05
```

```
No. Observations:      14448    AIC:
```

```
3.507e+05
```

```
Df Residuals:          14427    BIC:
```

```
3.509e+05
```

```
Df Model:              20
```

```
Covariance Type:      nonrobust
```

```
=====
=====
```

```
coef    std err          t    P>|t|
```

```
[0.025    0.975]
```

```
-----
```

```
const              -600.3963    65.349    -9.187    0.000
```

```
-728.489    -472.303
```

```
Median_Income      -3.018e+04    644.097    -46.849    0.000
```

```
-3.14e+04    -2.89e+04
```

```
Median_Age         -385.3993    37.300    -10.332    0.000
```

```
-458.512    -312.287
```

```
Tot_Rooms          9.559e+04    8545.582    11.185    0.000
```

```
7.88e+04    1.12e+05
```

|                          |            |                   |         |       |
|--------------------------|------------|-------------------|---------|-------|
| Tot_Bedrooms             | -7.085e+04 | 8646.634          | -8.194  | 0.000 |
| -8.78e+04                | -5.39e+04  |                   |         |       |
| Population               | -5.528e+04 | 1518.440          | -36.407 | 0.000 |
| -5.83e+04                | -5.23e+04  |                   |         |       |
| Households               | 3.36e+04   | 3029.159          | 11.094  | 0.000 |
| 2.77e+04                 | 3.95e+04   |                   |         |       |
| Latitude                 | -7.75e+04  | 8528.332          | -9.087  | 0.000 |
| -9.42e+04                | -6.08e+04  |                   |         |       |
| Longitude                | -2.464e+04 | 2380.125          | -10.353 | 0.000 |
| -2.93e+04                | -2e+04     |                   |         |       |
| Distance_to_coast        | -1.169e+04 | 822.861           | -14.207 | 0.000 |
| -1.33e+04                | -1.01e+04  |                   |         |       |
| Distance_to_LA           | -2.4097    | 0.821             | -2.936  | 0.003 |
| -4.018                   | -0.801     |                   |         |       |
| Distance_to_SanDiego     | -1.8097    | 0.478             | -3.789  | 0.000 |
| -2.746                   | -0.873     |                   |         |       |
| Distance_to_SanJose      | 12.7554    | 4.398             | 2.900   | 0.004 |
| 4.134                    | 21.377     |                   |         |       |
| Distance_to_SanFrancisco | -4.0410    | 5.046             | -0.801  | 0.423 |
| -13.931                  | 5.849      |                   |         |       |
| Badroom_Ratio            | 1.529e+05  | 3.15e+04          | 4.849   | 0.000 |
| 9.11e+04                 | 2.15e+05   |                   |         |       |
| Value_Income             | 0.1252     | 0.001             | 122.801 | 0.000 |
| 0.123                    | 0.127      |                   |         |       |
| Latitude_LA              | 0.0623     | 0.025             | 2.490   | 0.013 |
| 0.013                    | 0.111      |                   |         |       |
| Latitude_SanDiego        | 0.0581     | 0.014             | 4.065   | 0.000 |
| 0.030                    | 0.086      |                   |         |       |
| Longitude_SanJose        | 0.1063     | 0.036             | 2.961   | 0.003 |
| 0.036                    | 0.177      |                   |         |       |
| Longitude_SanFrancisco   | -0.0332    | 0.041             | -0.804  | 0.422 |
| -0.114                   | 0.048      |                   |         |       |
| LA_SanDiego              | -3.09e-07  | 1.54e-07          | -2.011  | 0.044 |
| -6.1e-07                 | -7.77e-09  |                   |         |       |
| SanJose_SanFrancisco     | -3.997e-07 | 1.05e-07          | -3.790  | 0.000 |
| -6.06e-07                | -1.93e-07  |                   |         |       |
| =====                    |            |                   |         |       |
| Omnibus:                 | 1955.250   | Durbin-Watson:    |         |       |
| 2.005                    |            |                   |         |       |
| Prob(Omnibus):           | 0.000      | Jarque-Bera (JB): |         |       |
| 26176.619                |            |                   |         |       |
| Skew:                    | -0.072     | Prob(JB):         |         |       |
| 0.00                     |            |                   |         |       |
| Kurtosis:                | 9.593      | Cond. No.         |         |       |
| 3.24e+15                 |            |                   |         |       |
| =====                    |            |                   |         |       |
| =====                    |            |                   |         |       |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.24e+15. This might indicate that there are strong multicollinearity or other numerical problems.

"""

```
test_data=x_test.join(y_test)
```

```
test_data
```

|       | Median_Income | Median_Age | Tot_Rooms | Tot_Bedrooms | Population |
|-------|---------------|------------|-----------|--------------|------------|
| \     |               |            |           |              |            |
| 9660  | 2.4519        | 22         | 7.661527  | 6.070738     | 6.721426   |
| 3033  | 3.2500        | 22         | 6.892642  | 5.176150     | 6.113682   |
| 3039  | 4.8266        | 13         | 8.418036  | 6.570883     | 7.682943   |
| 16586 | 3.4821        | 16         | 7.709757  | 6.129050     | 7.090077   |
| 16389 | 3.5735        | 25         | 7.584773  | 5.894403     | 6.943122   |
| ...   | ...           | ...        | ...       | ...          | ...        |
| 12755 | 3.0086        | 27         | 7.773174  | 6.287859     | 6.761573   |
| 16068 | 3.7361        | 48         | 7.687997  | 6.040255     | 6.948897   |
| 16681 | 4.4033        | 15         | 8.668884  | 6.829794     | 7.751475   |
| 15478 | 3.5839        | 5          | 8.466110  | 6.870053     | 7.910957   |
| 9510  | 3.8958        | 33         | 7.070724  | 5.356586     | 6.272877   |

|                  | Households | Latitude | Longitude | Distance_to_coast |
|------------------|------------|----------|-----------|-------------------|
| Distance_to_LA \ |            |          |           |                   |
| 9660             | 5.749393   | 41.31    | -121.18   | 12.393601         |
| 847236.324358    |            |          |           |                   |
| 3033             | 5.141664   | 35.39    | -119.11   | 11.671335         |
| 168509.993269    |            |          |           |                   |
| 3039             | 6.510258   | 35.37    | -119.12   | 11.650626         |
| 166991.195552    |            |          |           |                   |
| 16586            | 6.100319   | 37.75    | -121.44   | 10.458260         |
| 501862.376668    |            |          |           |                   |
| 16389            | 5.891644   | 38.05    | -121.25   | 10.090755         |
| 520163.946494    |            |          |           |                   |
| ...              | ...        | ...      | ...       | ...               |
| ...              |            |          |           |                   |
| 12755            | 6.115892   | 38.61    | -121.38   | 11.039695         |

|               |          |       |         |           |
|---------------|----------|-------|---------|-----------|
| 579364.293044 |          |       |         |           |
| 16068         | 5.940171 | 37.75 | -122.49 | 8.083626  |
| 561420.922514 |          |       |         |           |
| 16681         | 6.740519 | 35.13 | -120.56 | 9.100515  |
| 243545.558562 |          |       |         |           |
| 15478         | 6.809039 | 33.16 | -117.15 | 9.745623  |
| 141779.071052 |          |       |         |           |
| 9510          | 5.384495 | 39.13 | -123.23 | 10.652182 |
| 718742.255629 |          |       |         |           |

|       | ... | Distance_to_SanFrancisco | Bedroom_Ratio | Value_Income | \ |
|-------|-----|--------------------------|---------------|--------------|---|
| 9660  | ... | 407551.836660            | 0.203390      | 140984.25    |   |
| 3033  | ... | 397712.344960            | 0.178862      | 288925.00    |   |
| 3039  | ... | 398563.297627            | 0.157499      | 705648.92    |   |
| 16586 | ... | 87181.726267             | 0.205473      | 594046.26    |   |
| 16389 | ... | 108081.355786            | 0.184037      | 381649.80    |   |
| ...   | ... | ...                      | ...           | ...          |   |
| 12755 | ... | 130700.464955            | 0.226105      | 381791.34    |   |
| 16068 | ... | 5808.154770              | 0.192114      | 1196299.22   |   |
| 16681 | ... | 338278.378469            | 0.158817      | 1178323.08   |   |
| 15478 | ... | 701151.886604            | 0.202526      | 568048.15    |   |
| 9510  | ... | 166049.131134            | 0.179422      | 560995.20    |   |

|       | Latitude_LA  | Latitude_SanDiego | Longitude_SanJose | \ |
|-------|--------------|-------------------|-------------------|---|
| 9660  | 3.499933e+07 | 4.212630e+07      | -5.406774e+07     |   |
| 3033  | 5.963569e+06 | 1.229253e+07      | -3.930244e+07     |   |
| 3039  | 5.906479e+06 | 1.223568e+07      | -3.940261e+07     |   |
| 16586 | 1.894530e+07 | 2.571797e+07      | -7.406693e+06     |   |
| 16389 | 1.979224e+07 | 2.660674e+07      | -1.182640e+07     |   |
| ...   | ...          | ...               | ...               |   |
| 12755 | 2.236926e+07 | 2.926442e+07      | -1.804714e+07     |   |
| 16068 | 2.119364e+07 | 2.792438e+07      | -8.569656e+06     |   |
| 16681 | 8.555755e+06 | 1.450052e+07      | -3.289148e+07     |   |
| 15478 | 4.701394e+06 | 1.638460e+06      | -7.416967e+07     |   |
| 9510  | 2.812438e+07 | 3.513969e+07      | -2.848990e+07     |   |

|       | Longitude_SanFrancisco | LA_SanDiego  | SanJose_SanFrancisco | \ |
|-------|------------------------|--------------|----------------------|---|
| 9660  | -4.938713e+07          | 8.639781e+11 | 1.818403e+11         |   |
| 3033  | -4.737152e+07          | 5.853107e+10 | 1.312322e+11         |   |
| 3039  | -4.747686e+07          | 5.776789e+10 | 1.318371e+11         |   |
| 16586 | -1.058735e+07          | 3.419041e+11 | 5.317262e+09         |   |
| 16389 | -1.310486e+07          | 3.637284e+11 | 1.054197e+10         |   |
| ...   | ...                    | ...          | ...                  |   |
| 12755 | -1.586442e+07          | 4.391287e+11 | 1.943294e+10         |   |
| 16068 | -7.114409e+05          | 4.152935e+11 | 4.063506e+08         |   |
| 16681 | -4.078284e+07          | 1.005276e+11 | 9.228995e+10         |   |
| 15478 | -8.213994e+07          | 7.005408e+09 | 4.439112e+11         |   |
| 9510  | -2.046223e+07          | 6.454479e+11 | 3.838937e+10         |   |

Median\_House\_Value

|       |          |
|-------|----------|
| 9660  | 57500.0  |
| 3033  | 88900.0  |
| 3039  | 146200.0 |
| 16586 | 170600.0 |
| 16389 | 106800.0 |
| ...   | ...      |
| 12755 | 126900.0 |
| 16068 | 320200.0 |
| 16681 | 267600.0 |
| 15478 | 158500.0 |
| 9510  | 144000.0 |

[3096 rows x 22 columns]

OLS.score(x\_test,y\_test)

0.8589384353527918

ridge.score(x\_test,y\_test)

0.8589384353632443

lasso.score(x\_test,y\_test)

0.8566134110824559